

Assignment 2: Neural Networks

Due Date: Indicated on eLearning

Instructions

- There are two parts to this assignment. The first part requires you to solve some theoretical/numerical questions, and the second part requires you to code a neural network.
- For the programming part, please use parameters and not hard coded paths or values. All instructions for compiling and running your code must be placed in the README file.
- All work submitted must be your own. Do not copy from online sources. If you use any references, please list them.
- Please write the names and netids of all students on the front page.
- You are allowed to work in pairs i.e. a group of two students is allowed. Please write the names of the group members on the cover page.
- **You have a total of 4 free late days for the entire semester. You can use at most 2 days for any one assignment. After four days have been used up, there will be a penalty of 10% for each late day. The submission for this assignment will be closed 2 days after the due date.**
- Please ask all questions on Piazza, not via email.

1 Theoretical Part (40 points)

For the following, please show all steps of your derivation and list any assumptions that you make. You can submit typed or **legible** hand-written solutions. If the TA cannot read your handwriting, no credit will be given.

1.1 Revisiting Backpropagation Algorithm

In class we had derived the backpropagation algorithm for the case where each of the hidden and output layer neurons used the sigmoid activation function:

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$

Revise the backpropagation algorithm for the case where each hidden and output layer neuron uses the

a. tanh activation function

$$\tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

b. ReLu activation function:

$$\text{ReLU}(x) = \max(0, x)$$

Show all steps of your derivation and the final equation for output layer and hidden layers.

1.2 Gradient Descent

Derive a gradient descent training rule for a single unit neuron with output o , defined as:

$$o = w_0 + w_1(x_1 + x_1^2) + \dots + w_n(x_n + x_n^2)$$

where x_1, x_2, \dots, x_n are the inputs, w_1, w_2, \dots, w_n are the corresponding weights, and w_0 is the bias weight. You can assume an identity activation function i.e. $f(x) = x$. Show all steps of your derivation and the final result for weight update. You can assume a learning rate of η .

1.3 Comparing Activation Function

Consider a neural net with 2 input layer neurons, one hidden layer with 2 neurons, and 1 output layer neuron as shown in Figure 1. Assume that the input layer uses the identity activation function i.e. $f(x) = x$, and each of the hidden layers and output layer use an activation function $h(x)$. The weights of each of

the connections are marked in the figure.




Figure 1: A neural net with 1 hidden layer having 2 neurons

- Write down the output of the neural net y_5 in terms of weights, inputs, and a general activation function $h(x)$.
- Now suppose we use vector notation, with symbols defined as below:

$$X = \begin{pmatrix} x_1 \\ x_2 \end{pmatrix}$$

$$W^{(1)} = \begin{pmatrix} w_{3,1} & w_{3,2} \\ w_{4,1} & w_{4,2} \end{pmatrix}$$

$$W^{(2)} = (w_{5,3} \quad w_{5,4})$$

Write down the output of the neural net in vector format using above vectors.

- Now suppose that you have two choices for activation function $h(x)$, as shown below:

Sigmoid:

$$h_s(x) = \frac{1}{1 + e^{-x}}$$

Tanh:

$$h_t(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

Show that neural nets created using the above two activation functions can generate the same function.

Hint: First compute the relationship between $h_s(x)$ and $h_t(x)$ and then show that the output functions are same, with the parameters differing only by linear transformations and constants.

1.4 Gradient Descent with a Weight Penalty

Go through Chapter 4 of Tom Mitchell's Machine Learning textbook, which is available at the following link:

<https://users.cs.northwestern.edu/~pardo/courses/eecs349/readings/chapter4-ml.pdf>

Solve question 4.10 from the book. Show all steps of derivation and clearly state the final update rule for output as well as hidden layers.

2 Programming Part (60 points)

In this part, you will code a neural network (NN) having at least one hidden layers, besides the input and output layers. You are required to pre-process the data and then run the processed data through your neural net. Below are the requirements and suggested steps of the program

- The programming language for this assignment will be Python 3.x
- **You cannot use any libraries for neural net creation.** You are free to use any other libraries for data loading, pre-processing, splitting, model evaluation, plotting, etc.
- As the first step, pre-process and clean your dataset. There should be a method that does this.
- Split the pre-processed dataset into training and testing parts. You are free to choose any reasonable value for the train/test ratio, but be sure to mention it in the README file.
- Code a neural net having **at least one hidden layer**. You are free to select the number of neurons in each layer. Each neuron in the hidden and output layers should have a bias connection.
- You are required to add an optimizer on top of the basic backpropagation algorithm. This could be the one you selected in the previous assignment or a new one. Some good resources for gradient descent optimizers are:
<https://arxiv.org/pdf/1609.04747.pdf>
<https://ruder.io/optimizing-gradient-descent/>
<https://towardsdatascience.com/10-gradient-descent-optimisation-algorithms-86989510b5e9>
- Your code should be in the form of a Python class with methods like pre-process, train, test within the class. I leave the other details up to you.
- You are required to code three different activation functions:
 1. Sigmoid
 2. Tanh
 3. ReLu

The earlier part of this assignment may prove useful for this stage. The activation function should be a parameter in your code.

- Code a method for creating a neural net model from the training part of the dataset. Report the training accuracy.
- Apply the trained model on the test part of the dataset. Report the test accuracy.

- You have to tune model parameters like learning rate, activation functions, etc. Report your results in a tabular format, with a column indicating the parameters used, a column for training accuracy, and one for test accuracy.

Dataset

You can use **any one** dataset from the UCI ML repository:

<https://archive.ics.uci.edu/ml/datasets.php>

Note: If the above direct link does not work, you can just Google the UCI ML repository.

What to submit:

You need to submit the following for the programming part:

- Link to the dataset used. *Please do not include the data as part of your submission.*
- Your source code and a README file indicating how to run your code. Do not hardcode any paths to your local computer. It is fine to code any public paths, such as AWS S3.
- Output for your dataset summarized in a tabular format for different combination of parameters
- A brief report summarizing your results. For example, which activation function performed the best and why do you think so.
- Any assumptions that you made.